

# Assignment 1

Qiang Zhao (2814236)

November 2023

## 1 Answers

### question 1

a) The local derivatives of the softmax activation function contain two situations.

The first condition is as follows:

When  $i = j$ :

$$\begin{aligned}\frac{\partial y_i}{\partial o_j} &= \frac{\partial(\exp^{o_i} / \sum_j \exp^{o_j})}{\partial o_j} \\ &= \frac{\exp^{o_i} \sum_j \exp^{o_j} - \exp^{o_i} \exp^{o_i}}{(\sum_j \exp^{o_j})^2} \\ &= \frac{\exp^{o_i}}{\sum_j \exp^{o_j}} \left( 1 - \frac{\exp^{o_i}}{\sum_j \exp^{o_j}} \right) \\ &= y_i(1 - y_i)\end{aligned}$$

The second condition is as follows:

When  $i \neq j$ :

$$\begin{aligned}\frac{\partial y_i}{\partial o_j} &= \frac{\partial(\exp^{o_i} / \sum_j \exp^{o_j})}{\partial o_j} \\ &= -\frac{\exp^{o_i} \exp^{o_j}}{(\sum_j \exp^{o_j})^2} \\ &= -\frac{\exp^{o_i}}{\sum_j \exp^{o_j}} \frac{\exp^{o_j}}{\sum_j \exp^{o_j}} \\ &= -y_i y_j\end{aligned}$$

b) The local derivatives of the log loss function is as follows.

If  $c=i$ :

$$\begin{aligned}\frac{\partial l}{\partial y_i} &= \frac{\partial(-\ln(y_c))}{\partial y_i} \\ &= -\frac{1}{y_c}\end{aligned}$$

Otherwise:

$$\frac{\partial l}{\partial y_i} = 0$$

### question 2

If we assign the label as  $a_j$ , the local derivatives of the input of the softmax is as follows:

$$\begin{aligned}\frac{\partial l}{\partial o_i} &= \frac{\partial(-\sum_j a_j \ln(y_i))}{\partial o_i} \\&= -\sum_j a_j \frac{1}{y_i} \frac{\partial y_i}{\partial o_i} \\&= -a_j \frac{1}{y_i} y_i (1 - y_i) - \sum_{i \neq j} a_j \frac{1}{y_i} (-y_i y_j) \\&= a_j y_i - a_j + \sum_{i \neq j} a_j y_j \\&= \sum_j a_j y_j - a_j \\&= y_j \sum_j a_j - a_j \\&= y_j - a_j\end{aligned}$$

We do not need the step because we can get the same result using **the chain rule** ( $\frac{\partial l}{\partial o_i} = \frac{\partial l}{\partial y_i} \frac{\partial y_i}{\partial o_i}$ ).

### question 3

The relevant code was showing as follows:

```
#Setting up some predefined functions
def layer(input,output,weight,bias):
    n_unit = len(output)
    n_feature = len(input)
    for i in range(n_unit):
        for j in range(n_feature):
            output[i] += weight[j][i] * input[j]
        output[i] += bias[i]
    return output

def sigmoid(size,layer_1):
    s = [float(0) for i in range(size)]
    for i in range(len(s)):
        s[i] = math.exp(layer_1[i]) / (1+(math.exp(layer_1[i])))
    return s

def softmax(size,layer_2):
    layer_2_soft = np.zeros(len(layer_2)).tolist()
    layer_2_soft_deno = sum([math.exp(i) for i in layer_2])
    for j in range(len(layer_2)):
        layer_2_soft[j] = math.exp(layer_2[j]) / layer_2_soft_deno
    return layer_2_soft
```

```

def log_loss(pred,label):
    return -((label[0] * math.log(pred[0])) + (label[1] * math.log(pred[1])))

def one_hot_scalar(_,size):
    a = np.asarray([1 if i == _ else 0 for i in range(size)]).tolist()
    return a

#Define the class of scalar Multi-layer Perceptron
class scalarNN:
    def __init__(self,feature_size,layer_1_size,layer_2_size,fix_weight = True):
        self.fix = fix_weight
        self.l1_size = layer_1_size
        self.l2_size = layer_2_size

        if self.fix: #Fixed weight only for question 3
            self.weight_1 = [[1, 1, 1], [-1, -1, -1]]
            self.weight_2 = [[1, 1], [-1, -1], [-1, -1]]
        else: #Setting up weights from the Standard Normal Distribution
            self.weight_1 = [np.random.rand(self.l1_size).tolist() for i in range(feature_size)]
            self.weight_2 = [np.random.rand(self.l2_size).tolist() for i in range(self.l1_size)]

        self.bias_1 = [0 for i in range(self.l1_size)]
        self.bias_2 = [0 for i in range(self.l2_size)]

    def forward(self,input):
        layer_1_init = [float(0) for i in range(self.l1_size)]
        self.layer_1 = layer(input,layer_1_init,self.weight_1,self.bias_1)
        self.sigmoid_out = sigmoid(self.l1_size,self.layer_1)

        layer_2_init = [float(0) for i in range(self.l2_size)]
        self.layer_2 = layer(self.sigmoid_out,layer_2_init,self.weight_2,self.bias_2)
        self.softmax_out = softmax(self.l2_size,self.layer_2)
        return self.softmax_out

    def backward(self,x,pred,label,lr):
        layer_2_ld = [float(0) for i in range(self.l2_size)] #zero gradients before start
        for i in range(self.l2_size):
            layer_2_ld[i] = pred[i] - label[i]

        layer_1_sig_ld = [float(0) for i in range(self.l1_size)]
        for i in range(self.l1_size):
            for j in range(self.l2_size):
                layer_1_sig_ld[i] += layer_2_ld[j] * self.weight_2[i][j]

        layer_2_v_deri = [np.zeros(self.l2_size).tolist() for i in range(self.l1_size)]
        layer_2_c_deri = np.zeros(self.l2_size).tolist()
        for i in range(self.l2_size):
            for j in range(self.l1_size):
                layer_2_v_deri[j][i] = layer_2_ld[i] * self.sigmoid_out[j]
                self.weight_2[j][i] -= lr * layer_2_v_deri[j][i] #update weight v
            layer_2_c_deri[i] = layer_2_ld[i]
            self.bias_2[i] -= lr * layer_2_c_deri[i] #update bias c

```

```

layer_1_ld = np.zeros(self.l1_size).tolist()
for i in range(self.l1_size):
    layer_1_ld[i] = layer_1_sig_ld[i] * self.sigmoid_out[i] * (1 - self.sigmoid_out[i])

layer_1_w_der1 = [np.zeros(self.l1_size).tolist() for i in range(self.l2_size)]
layer_1_b_der1 = np.zeros(self.l1_size).tolist()
for i in range(self.l1_size):
    for j in range(self.l2_size):
        layer_1_w_der1[j][i] = layer_1_ld[i] * x[j]
        self.weight_1[j][i] -= lr * layer_1_w_der1[j][i] #update weight w
layer_1_b_der1[i] = layer_1_ld[i]
self.bias_1[i] -= lr * layer_1_b_der1[i] #update bias b

#Setting up the architecture of the network
feature_size = 2
layer_1_size = 3
layer_2_size = 2
lr = 0.1

#Network initialization
scalar_NN = scalarNN(feature_size,layer_1_size,layer_2_size,fix_weight=True)

#Prepare the input data
x = [1,-1]
label = [1,0]
pred = scalar_NN.forward(x)
loss = log_loss(pred, label)
scalar_NN.backward(x,pred,label,lr)

```

#### question 4

In question 4, we used the synthetic data from the following pre-defined function:

<https://gist.github.com/pbloem/bd8348d58251872d9ca10de4816945e4>

The training loop contained **100** epochs, each with **60,000** iterations; the parameters were updated for each instance. We calculated the average loss for each epoch. The relevant code was as follows:

```

#Load the synth data and normalizaiton
(xtrain, ytrain), (xval, yval), num_cls = load_synth()
data = xtrain
min_vals = np.min(data, axis=0)
max_vals = np.max(data, axis=0)
normal_data = (data - min_vals) / (max_vals - min_vals)

#Setting up the architecture of the network
feature_size = 2
layer_1_size = 3
layer_2_size = 2
lr = 0.1
iteration = normal_data.shape[0]
epoch = 100

#Network initialization
scalar_NN = scalarNN(layer_1_size,layer_2_size,fix_weight=False) #Close the fixed weight

#Feed the synthetic data to the network

```

```

y_loss = []
for k in range(epoch):
    total_loss = 0
    for i in range(1,iteration):
        x = list(i for i in list(normal_data[i-1]))
        label = one_hot_scalar(ytrain[i-1],2)
        print("epoch:",k+1,"round:",i,"/",normal_data.shape[0])
        pred = scalar_NN.forward(x)
        loss = log_loss(pred,label)
        total_loss += loss
        scalar_NN.backward(x,pred,label,lr)
    y_loss.append(total_loss / normal_data.shape[0])

```

*#Plot the lossing curve (The plot\_loss\_4 function was provided in the appendix part)*  
plot\_loss\_4(y\_loss,epoch,"Loss curve","Epoch","Average Loss over past one epoch")

From Figure 1, we can see the loss drops as training progresses, and it reaches convergence after round the 20th epoch.




Figure 1: The loss curve of Synthetic Data

## question 5 and 6

The code below demonstrates the process of implementing a **vectorized neural network**:

```

#Setting up some predefined functions
def log_loss_tensor(pred,label):
    pred = np.clip(pred, 1e-15, 1 - 1e-15)
    loss = - np.mean(np.log(pred)*label)
    return loss
def one_hot_vec(_,size):
    a = np.asarray([1 if i == _ else 0 for i in range(size) ]).reshape(size,1)
    return a

class tensorNN:
    def __init__(self,feature_size,layer_1_size,layer_2_size,batch_size):
        self.feature_s = feature_size #784

```

```

self.l1_s = layer_1_size #300
self.l2_s = layer_2_size #10
self.batch = batch_size

self.w = np.expand_dims(np.random.rand(self.l1_s) * 0.01,axis=0)
self.w = np.repeat(self.w, self.feature_s, axis=0).reshape(self.l1_s, self.feature_s)
self.b = np.array([0 for i in range(self.l1_s)]).astype('float64')

self.v = np.expand_dims(np.random.rand(self.l2_s) * 0.01,axis=0)
self.v = np.repeat(self.v, self.l1_s, axis=0).reshape(self.l2_s, self.l1_s)
self.c = np.array([0 for i in range(self.l2_s)]).astype('float64')

def forward(self,x):
    x = x.reshape(self.batch,self.feature_s)
    self.k1 = np.matmul(x,self.w.T) + self.b
    self.sigmoid_output = np.exp(self.k1) / (1+np.exp(self.k1))
    self.k2 = np.matmul(self.sigmoid_output,self.v.T) + self.c
    self.softmax_output = np.exp(self.k2) / (np.exp(self.k2).sum(axis=-1, keepdims=True))
    return self.softmax_output

def backward(self,x,label,lr):
    x = x.reshape(self.batch, self.feature_s)
    l2_ld = self.softmax_output - label
    l1_sig_ld = np.matmul(l2_ld, self.v)
    l2_v_der = np.matmul(l2_ld.T, self.sigmoid_output) / label.shape[0]
    l2_c_der = np.sum(l2_ld, axis=0) / label.shape[0]

    s_pr = self.sigmoid_output * (1 - self.sigmoid_output)
    l1_ld = s_pr * l1_sig_ld
    l1_w_der = np.matmul(l1_ld.T,x) / label.shape[0]
    l1_b_der = np.sum(l1_ld, axis=0) / label.shape[0]

    self.w -= lr * l1_w_der
    self.b -= lr * l1_b_der
    self.v -= lr * l2_v_der
    self.c -= lr * l2_c_der

#Setting up the shape of vectorized neural network
feature_size = 784
layer_1_size = 300
layer_2_size = 10
batch_size = 32
lr = 0.03

#Load the MNIST data
(xtrain, ytrain), (xval, yval), num_cls = load_mnist(final=False, flatten=True)
#Initialization of the Network
tensor_NN = tensorNN(feature_size,layer_1_size,layer_2_size,batch_size)
axis,yaxis,accuracy = [],[],0
rounds = int(xtrain.shape[0]/batch_size)
for r in range(rounds):
    start = (r) * batch_size
    end = (r + 1) * batch_size
    x_flatten = xtrain[start:end].reshape(batch_size,feature_size).astype('float32')

```

```

x_flatten /= 255 #Normalization of MNIST DATA
label = list(ytrain[start:end])
label_m = np.empty((batch_size, layer_2_size))
for i in range(len(label)):
    label_inter = one_hot_vec(label[i], layer_2_size)
    label_m[i, :] = label_inter.squeeze()
pred = tensor_NN.forward(x_flatten)
loss = log_loss_tensor(pred,label_m)
tensor_NN.backward(x_flatten, label_m, lr)
pred_max = np.argmax(pred, axis=1)
label_max = np.argmax(label_m, axis=1)
accuracy += np.mean(pred_max == label_max) #record accuracy for each batch
yaxis.append(loss)
xaxis.append(end)

accuracy = round(accuracy/rounds,3) #calculate average accuracy
#Plot the lossing curve (The plot_loss_5 function was provided in the appendix part)
plot_loss_5("Loss curve","Number of training samples",
"Average Loss over past one batch of samples (n=32)",accuracy,"05-one_epoch.png")

```

Because questions 5 and 6 both require the vectorized network, we combine the questions together. When implementing the network for one sample, you only need to set the parameter batch size to 1. The following analysis showed the result of batched gradient descent:

In tests on the **MNIST dataset**, we conducted the training process with a batch size of **32** and a learning rate of **0.03**. **Figure 2** shows the loss curve and average accuracy after just **one epoch**. **Figure 3** displays the loss curve and average accuracy after **100 epochs**. The above codes present the training process during one epoch. The code of 100 epochs was shown in the appendix part.




Figure 2: The loss curve of MNIST data after one epoch(left) and 100 epochs(right)

### question 7

1) The parameter used in this question is as follows:

The batch size is 32, and the learning rate is 0.03, as before.

**Figure 3** indicates the difference between the loss curves for the training and validation data, which lies in the fact that the training data starts with a lower loss and converges more quickly, achieving convergence after the fifth epoch. In contrast, the validation data begins with a higher initial loss and converges more slowly during the 5 epochs.

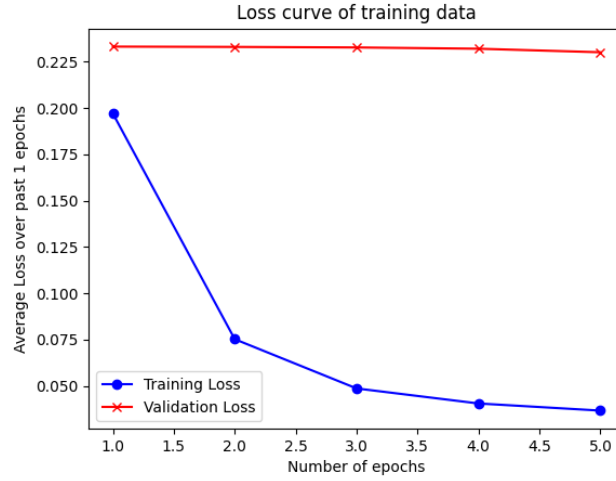


Figure 3: The loss curve of training data and validation data during 5 epochs

2)The parameter used in this question is as follows:

The batch size is **32**, and the learning rate is **0.03**, as before.




Figure 4: The loss curve of training data from 10 different initializations

From **Figure 4**, we can see that under the random initialization parameter (Implemented **10 times** in this question), as indicated by the SD (Standard Deviation) results, the loss curve fluctuates less with a smaller number of training samples. As training progresses, the SD gradually increases, leading to more significant fluctuations in the loss.

3) We employed two types of SGD approaches for comparison: stochastic gradient descent (**SGD**; batch size = 1) and mini-batch stochastic gradient descent (**mini-batch SGD**; batch size = 5,16,32). The comparison of the four sets of data in **Figure 5** shows that as the learning rate increases, the model enters the convergence process more rapidly. Interestingly, a higher batch size does not lead to a faster convergence for the same learning rate. The gradient of the smaller batch size groups starts to drop earlier than that of the batch size 32 group.



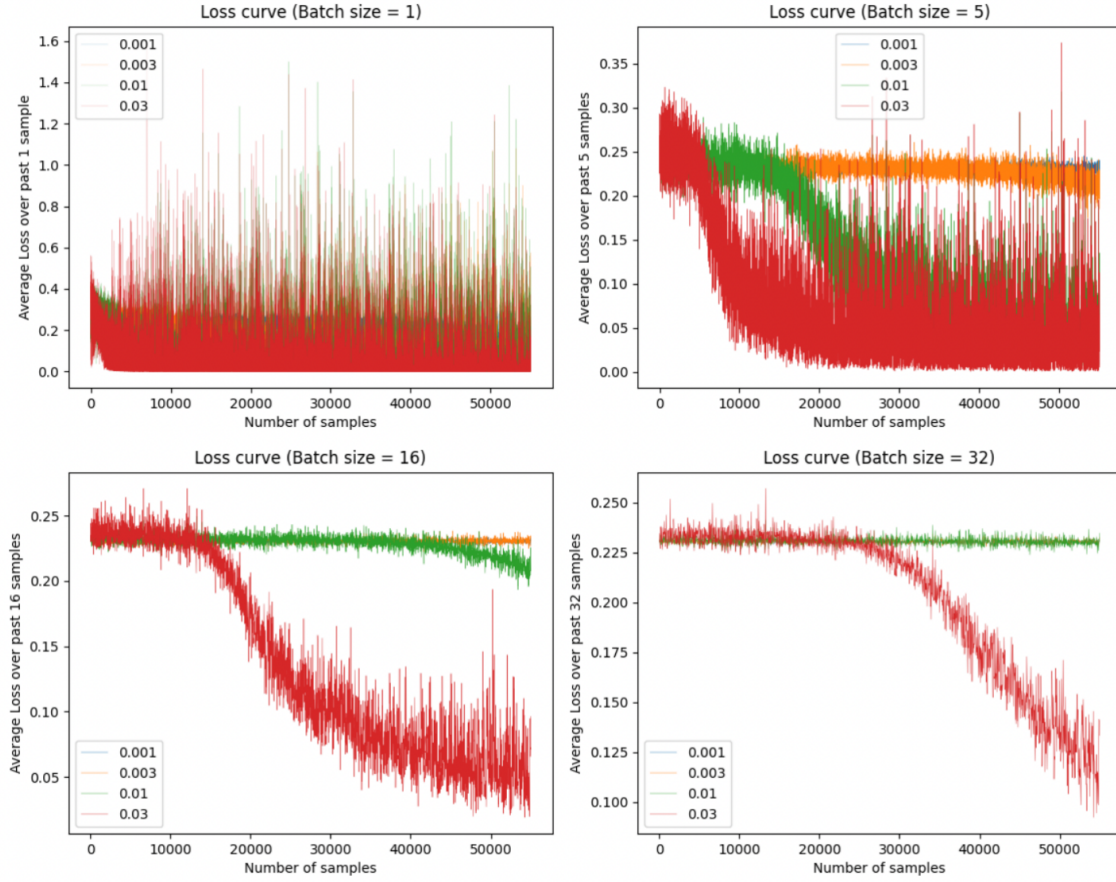


Figure 5: The loss curve training data with different learning rates and batch size

4) Based on the above analysis, we selected the following combination of hyperparameters:

a learning rate of **0.03** and a batch size of **16**.

We initially conducted five epochs on the complete MNIST training dataset. As shown in the left graph of **Figure 6**, the loss gradually decreased, and the accuracy steadily increased throughout training, reaching **0.916** in the fifth epoch. Subsequently, we performed a performance test for one epoch on the complete MNIST validation set. As seen in the right graph of **Figure 6**, the loss level remained stable, and the accuracy for this epoch reached **0.925**, which is higher than that of the training set. This indicates that our model has **good generalization ability** and did not exhibit overfitting.

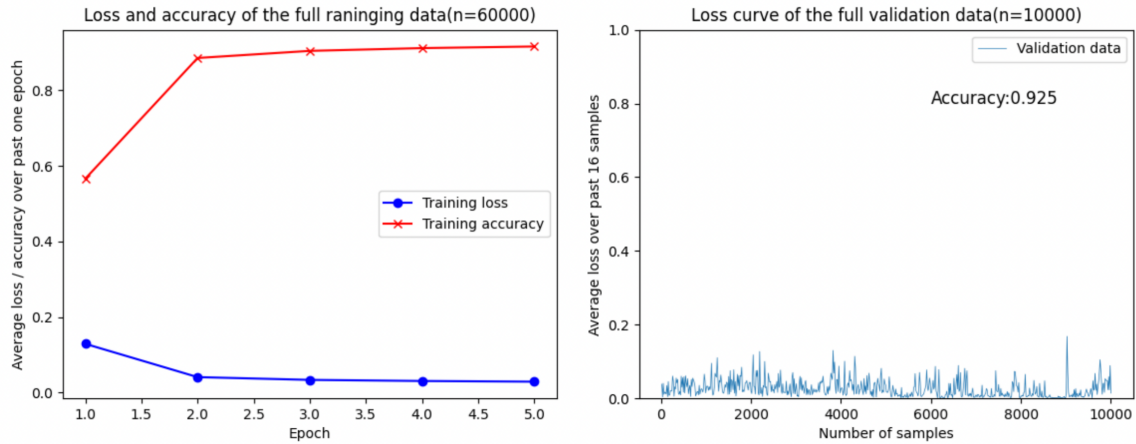


Figure 6: The performance of the final set of hyperparameters on the full training(left) and the full validation data(right)

## A Appendix

Here are some relevant codes for different questions. Some of the code was pasted with line breaks done manually to account for page margins.

### Question 4

The plot function for Figure 1

```
def plot_loss_4(loss_input,number_epoch,title,xtitle,ytitle):

    xaxis = list(range(1,number_epoch+1))
    yaxis = loss_input
    plt.plot(xaxis,yaxis,label='Synthetic Data')
    y_lim_max = max(yaxis) + 0.03
    y_lim_max = 1 if y_lim_max > 1 else y_lim_max

    plt.ylim(0, y_lim_max)
    plt.legend()
    plt.title(title)
    plt.xlabel(xtitle)
    plt.ylabel(ytitle)
    plt.savefig("04.png")
    plt.show()
```

### Question 5 and 6

The training and plot code for 100 epochs of MNIST data

```
#Setting the shape of the network
feature_size = 784
layer_1_size = 300
layer_2_size = 10
batch_size = 32
lr = 0.03
epoch = 100
```

```

#Load the MNIST data
(xtrain, ytrain), (xval, yval), num_cls = load_mnist(final=False, flatten=True)

#Initialization
tensor_NN = tensorNN(feature_size,layer_1_size,layer_2_size,batch_size)

axis,yaxis = [],[]
rounds = int(xtrain.shape[0]/batch_size)
y_loss = []
accuracy = 0

#Start to training the model
for k in range(epoch):
    total_loss = 0
    for r in range(rounds):
        start = (r) * batch_size
        end = (r + 1) * batch_size
        x_flatten = xtrain[start:end].reshape(batch_size,feature_size).astype('float32')
        x_flatten /= 255 #Normalization of MNIST DATA

        label = list(ytrain[start:end])
        label_m = np.empty((batch_size, layer_2_size))
        for i in range(len(label)):
            label_inter = one_hot_vec(label[i], layer_2_size)
            label_m[i, :] = label_inter.squeeze()

        pred = tensor_NN.forward(x_flatten)
        loss = log_loss_tensor(pred,label_m)
        total_loss += loss

        pred_max = np.argmax(pred, axis=1)
        label_max = np.argmax(label_m, axis=1)

        accuracy += np.mean(pred_max == label_max) #record accuracy for each batch
        tensor_NN.backward(x_flatten,pred,label_m,lr)
        print("epoch:", k + 1, "round:", r, "/", xtrain.shape[0])
    y_loss.append(total_loss / xtrain.shape[0])

accuracy = round(accuracy/(rounds*epoch),3) #calculate average accuracy

#Define the plot function
def plot_loss_5_epoch(y_loss,number_epoch,title,xtitle,ytitle,accuracy,output):

    xaxis = list(range(1, number_epoch + 1))
    yaxis = y_loss
    xtitle = xtitle
    ytitle = ytitle
    title = title

    a = "Average accuracy:" + str(accuracy)
    plt.text(max(xaxis)*0.1, max(yaxis)*0.2, a, fontsize=12)

    plt.plot(xaxis,yaxis,label='MNIST Data')
    y_lim_max = max(yaxis)*1.1

```

```

y_lim_max = 1 if y_lim_max > 1 else y_lim_max

plt.ylim(0, y_lim_max)

plt.title(title)
plt.xlabel(xtitle)
plt.ylabel(ytitle)
plt.savefig(output)
plt.legend()
plt.show()

#Plot the loss curve of Figure 2(right)
plot_loss_5_epoch(y_loss,epoch,"Loss curve","Epoch","Average Loss over past one epoch",accuracy,
"05-100_epoch.png")

```

### Question 7-1

The training and plot code for question 7-1:

```

#Setting the shape of the network
feature_size = 784
layer_1_size = 300
layer_2_size = 10
batch_size = 32
lr = 0.03

#Training_data
tensor_NN = tensorNN(feature_size,layer_1_size,layer_2_size,batch_size)
train = []
rounds = int(xtrain.shape[0]/batch_size)

for i in range(5):
    print("epoch:",i)
    train_epoch = 0
    for r in range(rounds):
        start = (r) * batch_size
        end = (r+1) * batch_size
        x_flatten = xtrain[start:end].reshape(batch_size,feature_size).astype('float32')
        x_flatten /= 255

        label = list(ytrain[start:end])
        label_m = np.empty((batch_size, layer_2_size))
        for i in range(len(label)):
            label_inter = one_hot_vec(label[i], layer_2_size)
            label_m[i, :] = label_inter.squeeze()

        pred = tensor_NN.forward(x_flatten)
        loss = log_loss_tensor(pred,label_m)
        tensor_NN.backward(x_flatten,label_m,lr)
        train_epoch += loss

    train.append(train_epoch/rounds)

#Validation data

```

```

tensor_NN = tensorNN(feature_size,layer_1_size,layer_2_size,batch_size)
val = []
rounds = int(xval.shape[0]/batch_size)
for i in range(5):
    print("epoch:",i)
    test_epoch = 0
    for r in range(rounds):
        start = (r) * batch_size
        end = (r+1) * batch_size
        x_flatten = xval[start:end].reshape(batch_size,feature_size).astype('float32')
        x_flatten /= 255

        label = list(yval[start:end])
        label_m = np.empty((batch_size, layer_2_size))
        for i in range(len(label)):
            label_inter = one_hot_vec(label[i], layer_2_size)
            label_m[i, :] = label_inter.squeeze()

        pred = tensor_NN.forward(x_flatten)
        loss = log_loss_tensor(pred,label_m)
        tensor_NN.backward(x_flatten,label_m,lr)
        test_epoch += loss

    val.append(test_epoch/rounds)

#Define the plot function
def plot_epoch(train,val,title,xtitle,ytitle):
    epochs = [1, 2, 3, 4, 5]
    train_loss = train
    val_loss = val

    plt.plot(epochs, train_loss, color='blue', marker='o', label='Training Loss')
    plt.plot(epochs, val_loss, color='red', marker='x', label='Validation Loss')

    plt.title(title)
    plt.xlabel(xtitle)
    plt.ylabel(ytitle)
    plt.legend()
    plt.savefig('07-01.png')
    plt.show()

#Plot the loss curve of figure 3
plot_epoch(train,val,"Loss curve of training data", "Number of epochs", "Average Loss
over past 1 epochs")

```

## Question 7-2

The training and plot code for question 7-2:

```

#Setting the shape of the network
feature_size = 784
layer_1_size = 300
layer_2_size = 10
batch_size = 32

```

```

lr = 0.03

rounds = int(xtrain.shape[0]/batch_size)

total_loss_x = [[] for _ in range(10)]
total_loss_y = []

for k in range(10):
    print("epoch:",k)
    tensor_NN = tensorNN(feature_size,layer_1_size,layer_2_size,batch_size)
    for r in range(rounds):
        start = (r)*batch_size
        end = (r+1)*batch_size
        x_flatten = xtrain[start:end].reshape(batch_size,feature_size).astype('float32')
        x_flatten /= 255 #normalization

        label = list(ytrain[start:end])
        label_m = np.empty((batch_size, layer_2_size))
        for i in range(len(label)):
            label_inter = one_hot_vec(label[i], layer_2_size)
            label_m[i, :] = label_inter.squeeze()

        pred = tensor_NN.forward(x_flatten)
        loss = log_loss_tensor(pred,label_m)
        total_loss_x[k].append(loss)
        if k == 0 : total_loss_y.append(end)
        tensor_NN.backward(x_flatten,label_m,lr)

def plot_mean_sd(x,y,title,xtitle,ytitle):
    y = y[0:len(x[0])]
    #for i in range(len(x)):
    # plt.plot(y,x[i],color = "blue",linewidth=0.2)
    x = np.array(x)
    mean_x = np.mean(x, axis=0)
    sd_x = np.std(x, axis=0)

    plt.plot(y,mean_x, color="blue",lw=0.3,label="Average loss")
    plt.fill_between(y, mean_x - sd_x, mean_x + sd_x,color='yellow', alpha=0.6,
    label = "Standard deviation(+ 1SD)")
    plt.title(title)
    plt.xlabel(xtitle)
    plt.ylabel(ytitle)
    plt.legend()
    plt.grid()
    plt.savefig('07-02.png')
    plt.show()

#Plot the loss curve of figure 4
plot_mean_sd(total_loss_x,total_loss_y,"Loss curve","Number of samples","Average Loss
over past 32 samples")

```

### Question 7-3

The training and plot code for question 7-3:

```

#Setting the shape of the network
feature_size = 784
layer_1_size = 300
layer_2_size = 10

batch_size = 1 # or 5,16,32 [Change the value manually everytime]
lr_list = [0.001, 0.003, 0.01, 0.03]

total_loss_x = [[],[],[],[]]
total_loss_y = []

rounds = int(xtrain.shape[0]/batch_size)

#Iterate all the learning rate
for k in range(len(lr_list)):
    lr = lr_list[k]
    print("epoch:",k)
    print("lr:",lr)
    tensor_NN = tensorNN(feature_size, layer_1_size, layer_2_size, batch_size)
    for r in range(rounds):
        start = (r)*batch_size
        end = (r+1)*batch_size
        x_flatten = xtrain[start:end].reshape(batch_size,feature_size).astype('float32')
        x_flatten /= 255
        label = list(ytrain[start:end])
        label_m = np.empty((batch_size, 10))
        for i in range(len(label)):
            label_inter = one_hot_vec(label[i], 10)
            label_m[i, :] = label_inter.squeeze()
        pred = tensor_NN.forward(x_flatten)
        loss = log_loss_tensor(pred,label_m)
        total_loss_x[k].append(loss)
        if k == 0:
            total_loss_y.append(end)
        tensor_NN.backward(x_flatten,label_m,lr)

def plot_lr(x,y,lr_list,title,xtitle,ytitle,output):
    x = x
    y = y
    for i in range(len(lr_list)):
        plt.plot(y, x[i], label=lr_list[i],lw=0.1)
    plt.title(title)
    plt.xlabel(xtitle)
    plt.ylabel(ytitle)
    plt.legend()
    plt.savefig(output)
    plt.show()

#Plot the loss curve of figure 4
#batchsize=1
plot_lr(total_loss_x,total_loss_y,lr_list,"Loss curve (Batch size = 1)",
"Number of samples","Average Loss over past 1 sample","07-03-batchsize=1.png")

```

```

#batchsize=5
plot_lr(total_loss_x,total_loss_y,lr_list,"Loss curve (Batch size = 5)",
"Number of samples","Average Loss over past 5 samples","07-03-batchsize=5.png")

#batchsize=16
plot_lr(total_loss_x,total_loss_y,lr_list,"Loss curve (Batch size = 16)",
"Number of samples","Average Loss over past 16 samples","07-03-batchsize=16.png")

#batchsize = 32
plot_lr(total_loss_x,total_loss_y,lr_list,"Loss curve (Batch size = 32)",
"Number of samples","Average Loss over past 32 samples","07-03-batchsize=32.png")

```

#### Question 7-4

The training and plot code for question 7-4:

```

#Load the final version of MNIST data
(xtrain_final, ytrain_final), (xval_final, yval_final), num_cls = load_mnist(final=True,
flatten=True)

#Set the final hyperparameters
feature_size = 784
layer_1_size = 300
layer_2_size = 10
batch_size = 16
lr = 0.03

#Start to training the full training_data with 5 epochs

tensor_NN = tensorNN(feature_size,layer_1_size,layer_2_size,batch_size)
train_x = []
train_y = []
accuracy_all_epoch = []
loss_all_epoch = []
rounds = int(xtrain_final.shape[0]/batch_size)
epoch = 5
for k in range(epoch):
    print("epoch:",i)
    accuracy = 0
    loss_epoch = 0
    for r in range(rounds):
        start = (r) * batch_size
        end = (r+1) * batch_size
        x_flatten = xtrain_final[start:end].reshape(batch_size,feature_size).astype('float32')
        x_flatten /= 255

        label = list(ytrain_final[start:end])
        label_m = np.empty((batch_size, layer_2_size))
        for i in range(len(label)):
            label_inter = one_hot_vec(label[i], 10)
            label_m[i, :] = label_inter.squeeze()

        pred = tensor_NN.forward(x_flatten)
        loss = log_loss_tensor(pred,label_m)
        tensor_NN.backward(x_flatten,label_m,lr)

```



```

        train_x.append(end)
        train_y.append(loss)
        pred_max = np.argmax(pred, axis=1)
        label_max = np.argmax(label_m, axis=1)
        accuracy += np.mean(pred_max == label_max)
        loss_epoch += loss
        print("epoch",k+1,"sampling",end,"/",xtrain_final.shape[0])
    loss_all_epoch.append(loss_epoch / rounds)
    accuracy_all_epoch.append(accuracy / rounds)

#Continue to use this model to test the full validation data in one epoch
val_x = []
val_y = []
accuracy = 0
rounds = int(xval_final.shape[0]/batch_size)
for i in range(1):
    print("epoch:", i)
    for r in range(rounds):
        start = (r)*batch_size
        end = (r+1)*batch_size
        x_flatten = xval_final[start:end].reshape(batch_size,feature_size).astype('float32')
        x_flatten /= 255 #normalization

        label = list(yval_final[start:end])
        label_m = np.empty((batch_size, layer_2_size))
        for i in range(len(label)):
            label_inter = one_hot_vec(label[i], 10)
            label_m[i, :] = label_inter.squeeze()

        pred = tensor_NN.forward(x_flatten)
        loss = log_loss_tensor(pred,label_m)
        tensor_NN.backward(x_flatten,label_m,lr)
        val_x.append(end)
        val_y.append(loss)
        pred_max = np.argmax(pred, axis=1)
        label_max = np.argmax(label_m, axis=1)
        accuracy += np.mean(pred_max == label_max)
    accuracy /= rounds
accuracy = round(accuracy,3)

#Define the plot function
def plot_training_epoch(loss_input,acc_input,title,xtitle,ytitle):
    epochs = [1, 2, 3, 4, 5]

    plt.plot(epochs, loss_input, color='blue', marker='o', label='Training loss')
    plt.plot(epochs, acc_input, color='red', marker='x', label='Training accuracy')

    plt.title(title)
    plt.xlabel(xtitle)
    plt.ylabel(ytitle)
    plt.legend()
    plt.savefig('07-04-training.png')
    plt.show()

```

```

def plot_loss_final(x,y,title,xtitle,ytitle,acc,output):

    xaxis = x
    yaxis = y
    plt.plot(xaxis,yaxis,lw=0.5,label = "Validation data")
    y_lim_max = 1

    plt.ylim(0, y_lim_max)
    a = "Accuracy:" + str(acc)

    plt.text(6000, 0.8, a, fontsize=12)
    plt.title(title)
    plt.xlabel(xtitle)
    plt.ylabel(ytitle)
    plt.legend()
    output = output
    plt.savefig(output)
    plt.show()

#Plot the loss and accuracy curve of figure 6 (left side)
plot_training_epoch(loss_all_epoch,accuracy_all_epoch,
"Loss and accuracy of the full training data(n=60000)","Epoch",
"Average loss / accuracy over past one epoch")

#Plot the loss curve of figure 6 (right side)
plot_loss_final(val_x,val_y,"Loss curve of the full validation data(n=10000)",
"Number of samples","Average loss over past 16 samples",accuracy,"07-04.validation.png")

```